

ECE482 — Introduction to Operating Systems

Project 1

Manuel — UM-JI (Fall 2022)

Goals of the project

- Write a simple shell
- Run the shell in Linux
- Run the shell in Minix 3

1 A tough life...

As you are comfortably sleeping in your soft bed, your alarm clock loudly rings and wakes you up. Already 11 AM... It would be good to sleep a bit longer but you already hear your mum stepping in the room and widely opening the curtains. As you are blinded by the sudden burst of light you start grumbling, but she is not even paying attention to you and just says: "I told you to sleep instead of playing video games all night long! You've graduated two months ago and you haven't found job yet!" The head under your pillow to protect yourself from the light you simply ignore her and start smiling as you think of last year at the same period when you had to wake up early to attend ve477 lectures. But luckily now you are home and you can sleep, so why not enjoying it? As you start to fall asleep again, she admonishes you "Wake up lazy kid! Time for you to find a job!"

"But mum, I have already sent two CVs last week? What else do you want me to do?" you try to argue.

"What do I want you to do? Hum...I want you to start doing something with your life! If you don't want to find a job then you'll work here!" she orders you.

As she carries on you simply look at her disconcerted. "Yes, I found some exercises for you: you're going to implement a shell," and she punctuates with an ironic smile.

Totally speechless, you look at her wondering how come she knows what a shell is.

As she approaches from your bed she hands in a couple of pages to you just saying "Lets work begin! And be sure that I'll check on your regularly..."

While you regain your composure you manage to articulate "I'm already very busy playing video games!"

To what she firmly replies "Not anymore you'll be busy implementing a shell!" Then she quits the room, leaving you puzzled and bitter. Why has not Lemonion call you already for an interview, you had the perfect profile? Now you are here lying in your bed with a couple of pages entitled *The mumsh shell*.

2 The mumsh shell

The main task of a shell is to wait for some user input, parse it and execute a command requested by the user. It should also provide support for input/output redirection and pipelining from a program into another one.

A shell simply consists of a loop that parses the user input. When waiting the shell displays a prompt, here the shell should display "mumsh \$ ".

When a command is input by the user it should be launched in a new process and the shell should block, waiting for the command to end.

A command line is composed of a command followed by some arguments. Arguments are space separated. The shell should exit when the user inputs `exit`. In case of error, such as when a command that does not exist is input, the shell should output an error on the standard error output (e.g. `Error: no such file or directory`).

The shell can be tested by comparing its behaviour to the result of the same commands in the regular Linux shell (e.g. `sh`, `bash`, `zsh`). Note that those commands are only for testing purpose, therefore they are far from being optimal and do not encompass all the features that need to be implemented.

```
mumsh $ echo 123 | grep 2
mumsh $ echo 123 > 1.txt
mumsh $ echo 456 >> 1.txt
mumsh $ cat < 1.txt
mumsh $ cat < 1.txt | sort -R > 2.txt
```

The `mumsh` shell is expected to be running in both Linux and Minix 3.

Hints:

- Useful system calls: `fork()`, `execvp()`, `wait/waitpid()`, `dup2/dup()`, `pipe()` and `close()`.
- A command line is not expected to be longer than 1024 characters.
- The use of the command `system` or of `lex` and `yacc` is prohibited.

3 Mum's Grading policy

Your mother has decided that you should follow a total of thirteen requirements, some with sub-tasks. For each requirement the awarded marks are displayed in bold into square brackets.

Important notes:

- Mum considers requirements marked with a * as advanced.¹
- A 50% penalty will be applied if commands are not launched in a new process;
- Any work that is not pushed onto the git server will be ignored;
- An issue should be opened for each task, and there should be at least one commit to the git server per task completed (more are recommended);
- Any task that does not exit cleanly, has memory leaks, or undefined behaviors will receive a **−10%** deduction;
- The shell should only output normal characters (e.g. no escape characters such as `\033` to handle color);
- Prepare a professional README file showing a “build badge” as well as basic installation, running instructions, and the author information. More advanced, but still simple, documentation must be part of a Wiki page. Lack of documentation will lead to a 10% deduction.

In the following description “requirement x” stands for requirement x, including all its sub-tasks, if any. A requirement having dependencies is considered completed if and only if it is completed together with all its dependencies.

¹Advanced requirements should only be completed if you enrolled in ECE4821.

1. Write a working read-parse-execute loop and an `exit` command; **[7.5]**
2. Handle single commands without arguments (e.g. `ls`); **[7.5]**
3. Support commands with arguments (e.g. `apt-get update` or `pkgin update`); **[8]**
4. File I/O redirection: **[5+5+5+2]**
 - 4.1. Output redirection by overwriting a file (e.g. `echo 123 > 1.txt`);
 - 4.2. Output redirection by appending to a file (e.g. `echo 465 >> 1.txt`);
 - 4.3. Input redirection (e.g. `cat < 1.txt`);
 - 4.4. Combine 4.1 and 4.2 with 4.3;
5. Support for bash style redirection syntax (e.g. `cat < 1.txt 2.txt > 3.txt 4.txt`); **[10]**
6. Pipes: **[5+5+5+10]**
 - 6.1. Basic pipe support (e.g. `echo 123 | grep 1`);
 - 6.2. Run all 'stages' of piped process in parallel. (e.g. `yes ve482 | grep 482`);
 - *6.3. Extend 6.2 to support requirements 4. and 5. (e.g. `cat < 1.txt 2.txt | grep 1 > 3.txt`);
 - *6.4. Extend 6.3 to support arbitrarily deep "cascade pipes" (e.g. `echo 123 | grep 1 | grep 1 | grep 1`);

Note: the sub-processes must be reaped in order to be awarded the marks.
7. Support `CTRL-D` (similar to bash, when there is no/an unfinished command); **[5]**
8. Internal commands: **[6+6+6]**
 - 8.1. Implement `pwd` as a built-in command;
 - 8.2. Allow changing working directory using `cd`;
 - 8.3. Allow `pwd` to be piped or redirected as specified in requirement 4.;
9. Support `CTRL-C`: **[5+3+2+10]**
 - 9.1. Properly handle `CTRL-C` in the case of requirement 4.;
 - *9.2. Extend 9.1 to support subtasks 6.1 to 6.3;
 - 9.3. Extend 9.2 to support requirement 7., especially on an incomplete input;
 - *9.4. Extend 9.3 to support requirement 6.;
- *10. Support quotes: **[5+3+3+5]**
 - 10.1. Handle single and double quotes (e.g. `echo "de'f' ghi" '123"a"bc' a b c`);
 - 10.2. Extend 10.1 to support requirement 4. and subtasks 6.1 to 6.3;
 - 10.3. Extend 10.2 in the case of incomplete quotes (e.g. Input `echo "de`, hit enter and input `cd`);
 - 10.4. Extend 10.3 to support requirements 4. and 6., together with subtask 9.3;
- *11. Wait for the command to be completed when encountering `>`, `<`, or `|`: **[3+3]**
 - 11.1. Support requirements 3. and 4. together with subtasks 6.1 to 6.3;
 - 11.2. Extend 11.1 to support requirement 9.;
12. Handle errors for all supported features. **[10]**

Note: a list of test cases will be published at a later stage. Marks will be awarded based on the number of cases that are correctly handled, i.e. if only if:

- A precise error message is displayed (e.g. simply saying "error happened!" is not enough);

- The program continues executing normally after the error is identified and handled;

*13. A command ending with an `&` should be run in background. **[10+5]**

- 13.1. For any background job, the shell should print out the command line, prepended with the job ID (e.g. if the two lines `/bin/ls &` and `/bin/ls | cat &` are input the output could be the two lines `[1] /bin/ls &` and `[2] /bin/ls | cat &`);
- 13.2. Implement the command `jobs` which prints a list of background tasks together with their running states (e.g. in the previous case output the two lines `[1] done /bin/ls &` and `[2] running /bin/ls | cat &`);

4 Mum's schedule

As your mother wants to closely monitor you and ensure you do not play video games in secret instead of working on her shell she has decided to enforce a tight schedule:

Milestone 1: tasks 1 to 5 must be completed;

Milestone 2: tasks 6 to 9 must be completed;

Final shell: fully functional shell;

It seems your mother has thought of everything, she has even organised to setup an Online Judge to verify that you are respecting the milestones; To show you she is willing to be flexible she allows you to be reasonably late for the milestones as long the final submission is completed on time. More precisely `m1` can be submitted until the deadline of `m2`, and `m2` until the final deadline. Each missed deadline will incur a -30% deduction on the final project grade. Besides for `m1` and `m2` passing between 25% and 50% of the test cases will lead to -10% , and -20% if less than 25% of the test cases are passed.

Feeling resigned you decide to follow her advice, so you grab your computer and start working in your bed thinking it makes a very comfortable desk...

A few minutes later she appears at your door and tells you: "By the way, here are some more detailed specs for your shell and some hints on how you can test it!" Before you realise it you see a USB stick falling in front of you. After opening the device you find a file containing the following information.

A First document

This file contains the following information on how you can test your shell.

JOJ Online Judge

Important reminders regarding the Online Judge (OJ):

- Save all the files, without any folder structure, into a `.tar` archive;
- Strictly stick to the specifications;
- The OJ only checks the correctness of the code not its quality;
- Use the provided `Makefile` or `CMakeLists.txt` files;

A.1 Input

Mum will test your shell with a carefully coded driver, so the input format in the test cases will be taken as the driver's `stdin`.

There are three kinds of commands in the input:

- `WRITELN <line>`: Write `line` to your shell ended with a `\n`.
- `CTRL <C|D>`: Send `SIGINT` with `CTRL C` or `EOF` with `CTRL D`.
- `WAIT <T>`: Wait for `T` ms.

A.2 Output

All of output should be flushed if it wasn't flushed automatically by `\n`. If you don't do so, the order of the output of different processes can be arbitrary, so you will get wrong answer.

The newline character (`\n`), tab character (`\t`) and spaces will be considered as the same in the output.

Mum will redirect your `stdout` and `stderr` to the same file in the driver, so please do not print debug information to `stderr`. The reason why it is done this way it is that a shell should print prompts and error messages to `stderr`, though our project doesn't have this requirement. Some may follow this rule but some will not, so mum has to consider them together.

A.3 Examples

The examples will show input in the left and output in the right. Most of them are from milestone 1 and 2, and the real test cases before. These are published in order to test your shell functionalities rather than sticking to your carefulness on I/O formats.

B Second document

This file provides more detailed information on mum's requirements as well as clearer specification for each task.

B.1 Write a working read-parse-execute loop and an exit command; [7.5]

The loop time should be reasonable (<10ms), your overall performance will be tested in this part.

Hint: Don't forget to print `exit` when exiting your shell.

Example (Milestone 1 Case 1)

```
WRITELN exit
```

```
mumsh $ exit
```

B.2 Handle single commands without arguments; [7.5]

Hint: there are only three files in the working directory, they are `driver`, `mumsh`, `mumsh_mmeory_check`.

Example (Milestone 1 Case 2)

WRITELN ls	mumsh \$ driver
WRITELN exit	mumsh
	mumsh_memory_check
	mumsh \$ exit

B.3 Support commands with arguments; [8]

After seeing some writing programs like `cat`, `ls` themselves (totally wrong in a shell), mum has these examples with package managers now.

Example (Milestone 1 Case 5)

WRITELN ls -a	mumsh \$.
WRITELN exit	..
	driver
	mumsh
	mumsh_memory_check
	mumsh \$ exit

B.4 File I/O redirection: [5+5+5+2]

This task is based on Task 2 and 3.

Only one output redirection and one input redirection is allowed in a command. The test cases will strictly follow this rule.

B.4.1 Output redirection by overwriting a file [5]

Example (Milestone 1 Case 7)

WRITELN echo 321 > 1.txt	mumsh \$ mumsh \$ mumsh \$ 123
WRITELN echo 123 > 2.txt	mumsh \$ 321
WRITELN cat 2.txt	mumsh \$ 1.txt
WRITELN cat 1.txt	2.txt
WRITELN ls	driver
WRITELN exit	mumsh
	mumsh_memory_check
	mumsh \$ exit

B.4.2 Output redirection by appending to a file; [5]

Example (Milestone 1 Case 9)

```
WRITELN echo 321 >> 1.txt      mumsh $ mumsh $ 321
WRITELN cat 1.txt             mumsh $ 1.txt
WRITELN ls                    driver
WRITELN exit                  mumsh
                              mumsh_memory_check
                              mumsh $ exit
```

B.4.3 Input redirection; [5]

Example (Milestone 1 Case 11)

```
WRITELN echo 123 > 1.txt      mumsh $ mumsh $ 123
WRITELN cat < 1.txt           mumsh $ 1.txt
WRITELN ls                    driver
WRITELN exit                  mumsh
                              mumsh_memory_check
                              mumsh $ exit
```

B.5 Support for bash style redirection syntax; [10]

The bash style redirection syntax is very complex, You are only required to implement part of it. The following example actually shows all of the requirements:

- An arbitrary amount of space can exist between redirection symbols and arguments, starting from zero.
- The redirection symbol can take place anywhere in the command.

Example (Milestone 1 Case 18)

```
WRITELN echo 123 > 1.txt      mumsh $ mumsh $ mumsh $ mumsh $ mumsh $ 123
WRITELN echo 222 > 2.txt      mumsh $ 222
WRITELN echo 321 > 4.txt      mumsh $ 222
WRITELN <1.txt>3.txt cat 2.txt 4.txt 321
WRITELN cat 1.txt             mumsh $ 321
WRITELN cat 2.txt             mumsh $ 1.txt
WRITELN cat 3.txt             2.txt
WRITELN cat 4.txt             3.txt
WRITELN ls                    4.txt
WRITELN exit                  driver
                              mumsh
                              mumsh_memory_check
                              mumsh $ exit
```

B.6 Pipes: [5+5+5+10]

B.6.1 Basic pipe support; [5]

Example (Milestone 2 Case 1)

WRITELN ls -a grep mum	mumsh \$ mumsh
WRITELN exit	mumsh_memory_check
	mumsh \$ exit

B.6.2 Run all “stages” of piped process in parallel; [5]

Your total runtime is limited to 1s for the test.

Example (Milestone 2 Case 3)

WRITELN sleep 0.6 sleep 0.6	mumsh \$ mumsh \$ success
WRITELN echo success	mumsh \$ exit
WRITELN exit	

A tricky test of executing the shell itself.

Example (Milestone 2 Case 4)

WRITELN echo exit ./mumsh	mumsh \$ mumsh \$ exit
WRITELN exit	mumsh \$ exit

B.6.3 Extend 6.2 to support requirements 4. and 5. ; [5]

Input redirection can only be found in the first command, and output redirection can only be found in the last command, and each of them can occur only once. The test cases will strictly follow this rule.

Example (Milestone 2 Case 5)

WRITELN echo 123 > 1.txt	mumsh \$ mumsh \$ mumsh \$ mumsh \$ 123
WRITELN echo 321 > 2.txt	mumsh \$ 321
WRITELN cat < 1.txt cat >> 2.txt	123
WRITELN cat 1.txt	mumsh \$ 1.txt
WRITELN cat 2.txt	2.txt
WRITELN ls	driver
WRITELN exit	mumsh
	mumsh_memory_check
	mumsh \$ exit

B.6.4 Extend 6.3 to support arbitrarily deep “cascade pipes” ; [10]

Do you like the cats in the example? (Miao -,-)


```
Example (Milestone 2 Case 8)
```

WRITELN echo 123 > 1.txt	mumsh \$ mumsh \$ mumsh \$ mumsh \$ 123
WRITELN echo 321 > 2.txt	mumsh \$ 321
WRITELN cat < 1.txt cat cat cat cat cat cat cat cat cat cat cat cat cat	
WRITELN cat 1.txt	mumsh \$ 1.txt
WRITELN cat 2.txt	2.txt
WRITELN ls	driver
WRITELN exit	mumsh mumsh_memory_check mumsh \$ exit

```
Example (Milestone 2 Case 13)
```

WRITELN pwd > 1.txt	mumsh \$ mumsh \$ /out/package
WRITELN cat 1.txt	mumsh \$ 1.txt
WRITELN ls	driver
WRITELN exit	mumsh
	mumsh_memory_check
	mumsh \$ exit

You should NOT use `kill(0, signal)`, it may destroy the judge. You will be sent to the honor council if you intentionally use it.

B.9.1 Properly handle CTRL-C in the case of requirement 4. [5]

Example (Milestone 2 Case 15)

[illegible]

B.9.2 Extend 9.1 to support subtasks 6.1 to 6.3; [3]

Example (Milestone 2 Case 16)

```
WRITELN sleep 10 | echo 123          mumsh $ 123
CTRL C
WRITELN echo success                mumsh $ success
WRITELN exit                        mumsh $ exit
```

B.9.3 Extend 9.2 to support requirement 7., especially on an incomplete input; [2]

Example (Milestone 2 Case 17)

```
CTRL C                                mumsh $
CTRL C                                mumsh $
                                     mumsh $ exit
```

B.9.4 Extend 9.3 to support requirement 6.; [10]

Example (Milestone 2 Case 18)

```

WRITELN sleep 10 | sleep 10 | sleep 10 | sleep 10 | sleep 10 | sleep 10 | sleep 10 | sleep 10 | sleep
CTRL C
CTRL C
WRITELN echo success
CTRL C
WRITELN exit

```

B.10 Support quotes: [5+3+3+5]

B.10.1 Handle single and double quotes; [5]

Hint: the program name can also be quoted.

Example (Final Pretest Case 2)

```
WRITELN "echo" 'abc def' "xm"fxh"      mumsh $ abc def xmfxfh
WRITELN exit                             mumsh $ exit
```

B.10.2 Extend 10.1 to support requirement 4. and subtasks 6.1 to 6.3; [3]

Example (Final Pretest Case 3)

```
WRITELN echo 123 > 1.txt                mumsh $ mumsh $ mumsh $ 123
WRITELN "echo" "<1.'txt'" < 1.txt > "2.'txt'" mumsh $ <1.'txt'
WRITELN cat 1.txt                        mumsh $ 1.txt
WRITELN cat 2.txt                        2.txt
WRITELN ls                               driver
exit                                     mumsh
                                         mumsh_memory_check
                                         mumsh $ exit
```

Example (Final Pretest Case 4)

```
WRITELN "echo" "123" "|" | cat > 1.txt   mumsh $ mumsh $ 123 |
WRITELN cat 1.txt                        mumsh $ 1.txt
WRITELN ls                               driver
exit                                     mumsh
                                         mumsh_memory_check
                                         mumsh $ exit
```

B.10.3 Extend 10.2 in the case of incomplete quotes; [3]

This part is similar to some functions in task 11. Consider them together. You should output "> " after the user hits enter.

Example (Final Pretest Case 6)

```
WRITELN echo 123 > 1.txt                mumsh $ mumsh $ > mumsh $ 123
WRITELN echo "<1.'txt'"                mumsh $ <1.'txt'
WRITELN " < 1.txt >> 2.'txt'"          mumsh $ 1.txt
WRITELN cat 1.txt                        2.'txt'
WRITELN cat 2.'txt'"                    driver
WRITELN ls                               mumsh
WRITELN exit                             mumsh_memory_check
                                         mumsh $ exit
```

B.10.4 Extend 10.3 to support requirements 4. and 6., together with subtask 9.3; [5]

Usually you do not need to modify many things for this task if you design your shell well. It is not tested in the final pretest, but will be tested in the final shell.

B.11 Wait for the command to be completed when encountering >, <, or |: [3+3]

You should output "> " after the user hits enter.

B.11.1 Support requirements 3. and 4. together with subtasks 6.1 to 6.3; [3]

Hint: "»" does not need to be supported.

Example (Final Pretest Case 7)

WRITELN echo 123 >	mumsh \$ > mumsh \$ 123
WRITELN 1.txt	mumsh \$ 1.txt
WRITELN cat 1.txt	driver
WRITELN ls	mumsh
WRITELN exit	mumsh_memory_check
	mumsh \$ exit

Example (Final Pretest Case 9)

WRITELN echo 321 > 1.txt	mumsh \$ mumsh \$ > 321
WRITELN cat <	mumsh \$ 1.txt
WRITELN 1.txt	driver
WRITELN ls	mumsh
WRITELN exit	mumsh_memory_check
	mumsh \$ exit

Example (Final Pretest Case 11)

WRITELN echo 321	mumsh \$ > 321
WRITELN cat	mumsh \$ exit
WRITELN exit	

B.11.2 Extend 11.1 to support requirement 10.; [3]

Hint: Actually 10.3 already has this functionality, so you only need to extend to 10.1, 10.2, 10.4.

Example (Final Pretest Case 13)

WRITELN <code>echo</code> 123 > 1.txt	mumsh \$ mumsh \$ > > > mumsh \$ 123
WRITELN <code>"echo"</code> "<	mumsh \$ <
WRITELN <code>1.'</code> txt'> <	1.'
WRITELN 1.txt >	mumsh \$ 1.txt
WRITELN <code>"2."</code> txt"	2.txt
WRITELN cat 1.txt	driver
WRITELN cat 2.txt	mumsh
WRITELN ls	mumsh_memory_check
<code>exit</code>	mumsh \$ <code>exit</code>

B.12 Handle errors for all supported features. [10]

There are eight kinds of errors to be handled. In the test cases of other tasks, none of these errors will be included. You will always get a correct command then. At most one error will occur in a command, so you do not need to consider about the order of errors.

The output format of your errors should strictly follow the format listed below. It is very similar to the bash style error output format.

1. Non-existing program
 - input: `non-exist abc def`
 - input: `echo abc | non-exist`
 - output: `non-exist: command not found`
2. Non-existing file in input redirection
 - input: `cat < non-existing.txt`
 - output: `non-existing.txt: No such file or directory`
3. Failed to open file in output redirection
 - input: `echo abc > /dev/permission_denied`
 - output: `/dev/permission_denied: Permission denied`
4. Duplicated input redirection
 - input: `echo abc < 1.txt < 2.txt`
 - input: `echo abc | grep abc < 1.txt`
 - output: `error: duplicated input redirection`
5. Duplicated output redirection
 - input: `echo abc > 1.txt > 2.txt`
 - input: `echo abc > 1.txt >> 2.txt`
 - input: `echo abc > 1.txt | grep abc`
 - output: `error: duplicated output redirection`
6. Syntax Error
 - input: `echo abc > > > >`

- output: syntax error near unexpected token `>'
- input: `echo abc > < 1.txt`
- output: syntax error near unexpected token `<'
- input: `echo abc > | grep abc`
- output: syntax error near unexpected token `|'

7. Missing program

- input: `echo 123 | | grep 123`
- output: error: missing program

8. cd to non-existing directory

- input: `cd non-existing`
- output: non-existing: No such file or directory

B.13 A command ending with an & should be run in background. [10+5]

B.13.1 For any background job, the shell should print out the command line, prepended with the job ID [10]

Hints:

- Do not print the process ID.
- Do not print complete information or "mumsh \$" after the job is finished (like 'bash' or 'zsh').
- The job ID starts from 1.

Example (Final Pretest Case 23)

WRITELN /bin/ls &	mumsh \$ [1] /bin/ls &
WRITELN sleep 0.2	mumsh \$ driver
WRITELN /bin/ls cat &	mumsh
WRITELN sleep 0.2	mumsh_memory_check
WRITELN exit	mumsh \$ [2] /bin/ls cat &
	mumsh \$ driver
	mumsh
	mumsh_memory_check
	mumsh \$ exit

B.13.2 Implement the command jobs which prints a list of background tasks together with their running states [5]

Hint: Print all of done and running background processes in this shell.

Example (Final Pretest Case 24)

```
WRITELN /bin/ls &
WRITELN sleep 0.2
WRITELN sleep 0.2 &
WRITELN jobs
WRITELN sleep 0.4
WRITELN exit

mumsh $ [1] /bin/ls &
mumsh $ driver
mumsh
mumsh_memory_check
mumsh $ [2] sleep 0.2 &
mumsh $ [1] done /bin/ls &
[2] running sleep 0.2 &
mumsh $ mumsh $ exit
```